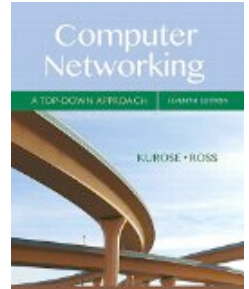


COMP 375: Lecture 24



- **News & Notes:**

- Spring Break... WOOOOOO!!!
- Special presentation at end of class
- Quiz #6 in class Wednesday, April 4
- Project #4 due Monday, April 16


- **Reading (Wed, April 4)**

- Sections 4.3.{0-3} (IPv4 Packets and Addressing, excluding 4.3.3 on DHCP)
- Note: Older edition chapters get out of sync at this point...

Sections 3.6 – 3.7

CONGESTION CONTROL

Sender's rate is a function of cwnd and RTT.



$$\text{rate} \approx \frac{\text{cwnd}}{\text{RTT}} \text{ bytes/sec}$$

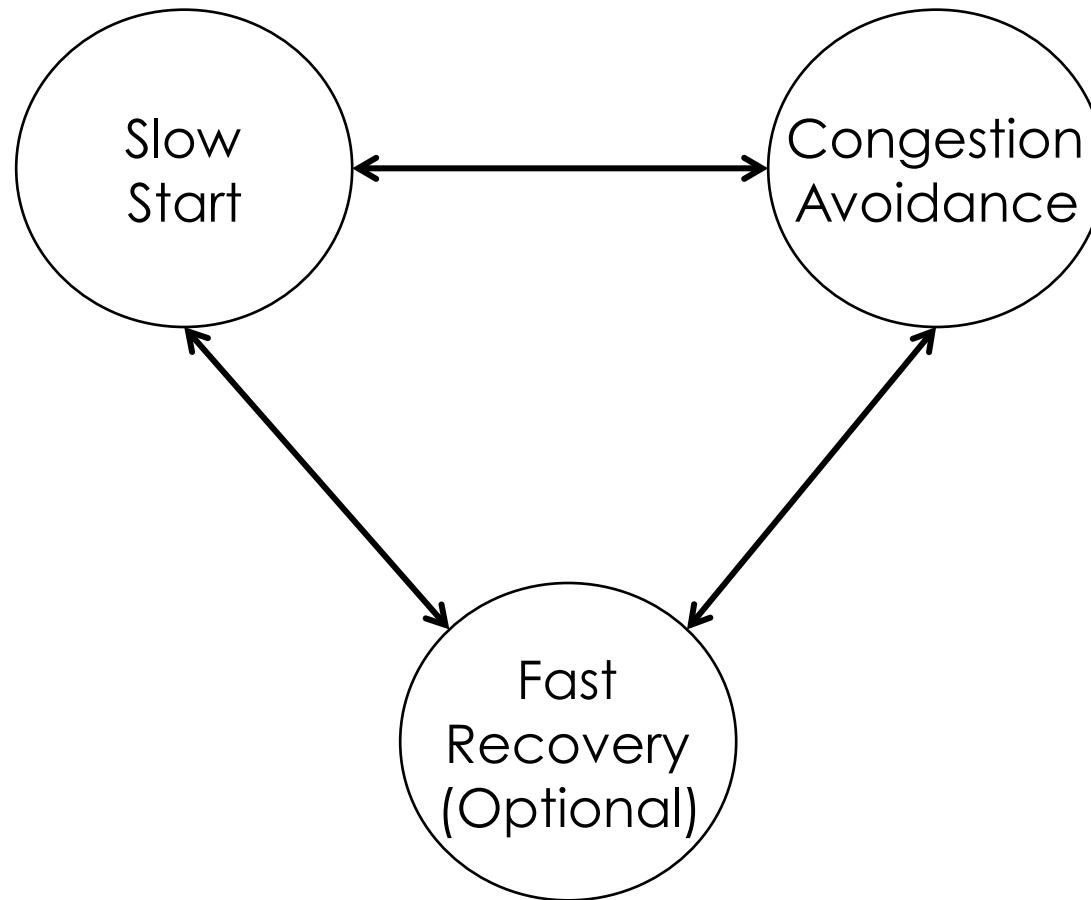
How should we set cwnd?

- A.** We should keep raising it until a “congestion event”, then back off slightly until we notice no more events.
- B.** We should raise it until a “congestion event”, then go back to 1 and start raising it again.
- C.** We should raise it until a “congestion event”, then go back to a median value and start raising it again.
- D.** We should send as fast as possible at all times.


What is a “congestion event”?

- | | |
|-----------|----------------------------|
| A. | A segment loss |
| B. | Receiving duplicate ACK(s) |
| C. | Timeout |
| D. | Exactly 2 of the above. |
| E. | A, B, and C |

TCP goes through phases that dictate how `cwnd` changes.



In the **slow start** state, we start slow but increase rate exponentially.




Eventually we will need to transition away from slow start.

- TCP leaves the slow start state for one of **two** reasons...

What are those reasons?

In the **congestion avoidance** state we increase rate linearly.



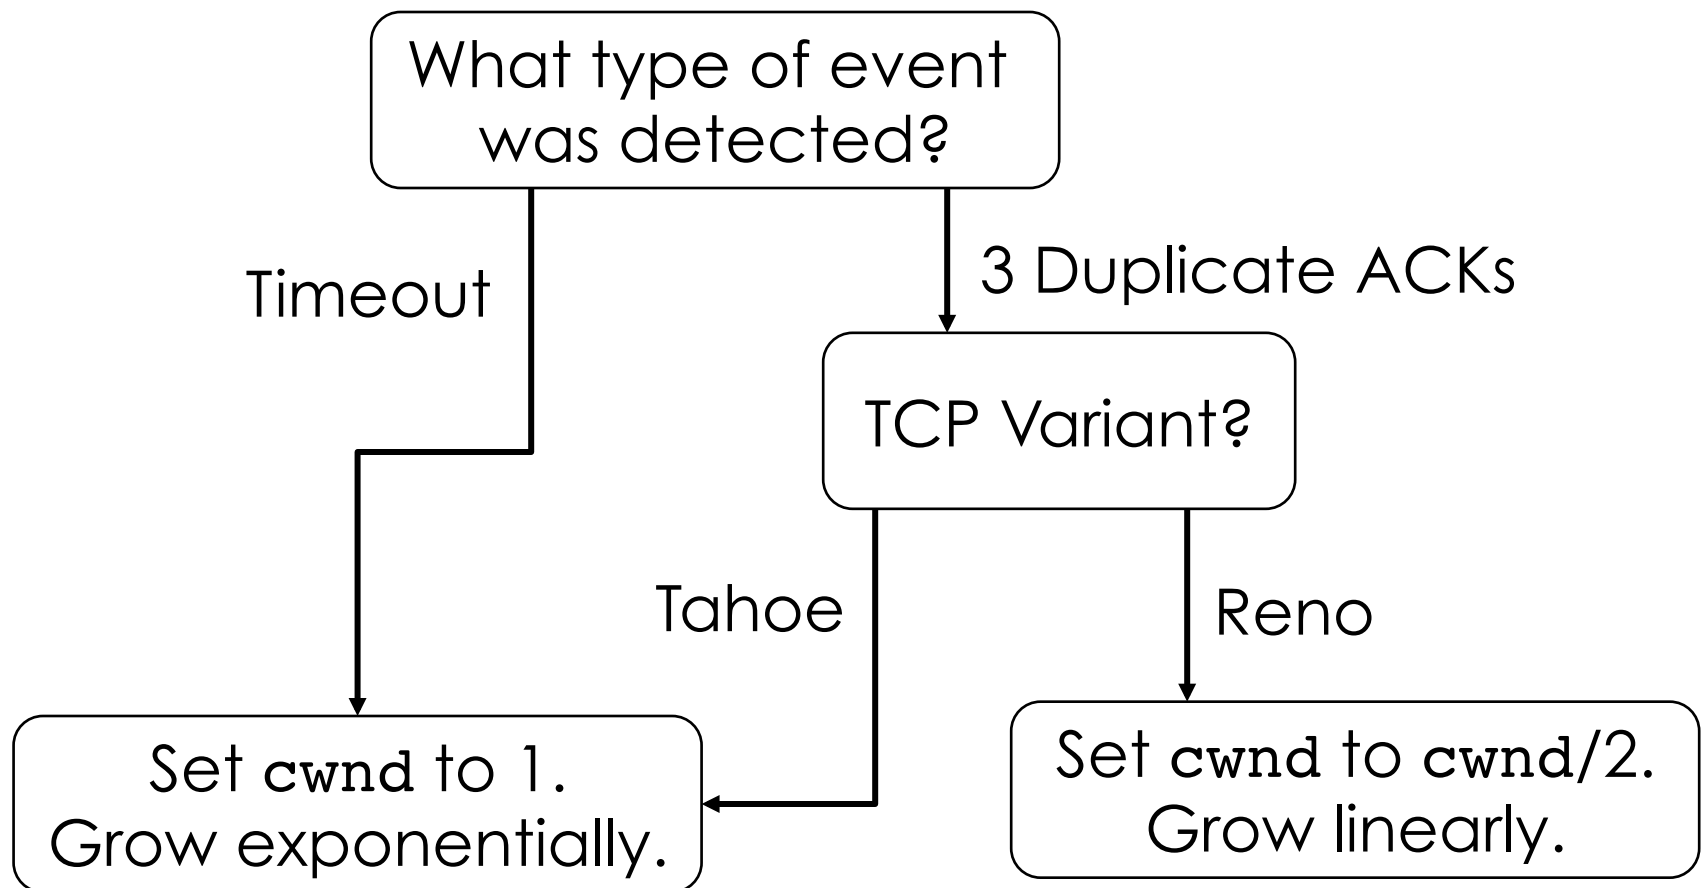
We can detect loss via three duplicate ACKs, or via a timeout.

How should we respond to these two types of loss detection?


- A.** Treat these events differently.
- B.** Treat these events the same.

Discuss: Which, if either, of these events are “worse.”


TCP variants react differently to loss.




`cwnd`'s sawtooth shape comes from AIMD
(Additive Increase, Multiplicative Decrease)



`ssthresh` dictates the transition between slow start and congestion avoidance.



Our goal is for K TCP connections to share bottleneck bandwidth equally (i.e. R/K).



AIMD helps ensure fairness in TCP.

